

DESK-AI: AN AUTONOMOUS, LOCAL-FIRST CONTINUOUS SCREEN & AUDIO CONTEXT ENGINE FOR AGENTIC COMPUTING

“FROM MULTIMODAL DESKTOP INGESTION TO LOCAL ZERO-KNOWLEDGE EPISODIC REASONING”



Presented by:

Sauhard Dubey (9923102092)

Anubhav Maurya (9923102097)

Mahatva Singh (9923102005)

Under the Supervision of:

Dr. Ruchi Singh

Department of ECE / CSE

Jaypee Institute of Information Technology, Noida, U.P., India

Table of Contents

Comprehensive Roadmap of the Capstone Project Presentation

PART I: RESEARCH & MOTIVATION

- Project Scope: Sovereign Ambient Intelligence
- Problem Formulation & End-Term Work Summary
- Literature Survey: Evolution of Desktop Memory
- Problem Statement: The Digital Context Bottleneck
- Introducing Desk-AI: The Local-First Paradigm
- System Architecture: Multi-Tier Modular Topology
- End-to-End Multimodal Ingestion Pipeline

PART II: SYSTEMS ENGINEERING & BENCHMARKS

- Deep Dive: Event-Driven Screen Ingestion & dHash
- Deep Dive: Dual-Stream Audio, Silero VAD & Local STT
- Deep Dive: Zero-Knowledge DESKVAULT1 Encryption
- Backend Data Architecture: SQLite WAL & FTS5 Indexing
- User Interface: Desktop DVR Timeline & Control Center
- Agent Interoperability: Model Context Protocol (MCP)
- Evolution: Mid-Term Prototype vs. End-Term System
- Experimental Benchmarks & Resource Profiling
- Why Commercial Systems Fail vs. Desk-AI Advantages
- Verification, Test Case Matrix & References

Scope of the Project: Sovereign Ambient Intelligence

Addressing the Universal Context Bottleneck Without Cloud Surveillance

- **Universal Cognitive Bottleneck:** Modern knowledge workers navigate dozens of applications daily (IDEs, terminal sessions, browsers, conference calls), absorbing >10 Mbps of multimodal information. Biological episodic memory decays rapidly, inducing severe context re-acquisition latency when recalling terminal commands, fleeting error logs, or meeting decisions.
- **Zero-Knowledge Sovereign Infrastructure:** Unlike commercial cloud-hosted lifeloggers (Otter, Limitless, Recall), Desk-AI is architected from the ground up as a 100% on-device, zero-telemetry engine. All audio waveforms, screen frames, and accessibility text reside exclusively in local storage, guaranteeing full compliance with enterprise security, GDPR, and HIPAA standards.
- **Cross-Agent Ecosystem Impact:** Scientific and software engineering workflows are inherently interdisciplinary. By standardizing continuous screen and audio context over the Model Context Protocol (MCP), Desk-AI bridges human workflow history directly into AI coding agents (Claude, Cursor, Codex), eliminating prompt fatigue and manual copy-pasting.

End-Term Work Summary: Problem Formulation & Its Solution

Overcoming Computational Overhead, OCR Inaccuracy, and Security Liabilities

• Problem Identification:

We identified critical systems bottlenecks in conventional continuous desktop memory implementations:

- Continuous fixed-framerate video capture (15–30 FPS) consumes prohibitive CPU cycles (20–40%) and degrades laptop battery life.
- Brute-force Optical Character Recognition (OCR) is computationally slow, lossy, and completely ignores semantic UI hierarchies and browser URLs.
- Cloud exfiltration and unencrypted local databases expose proprietary intellectual property, passwords, and private conversations.

• How We Solved the Above Problems:

1. Event-Driven Capture & Perceptual Diffing: Constructed OS event hooks (clicks, window focus, typing pauses) paired with a 64-bit difference hashing (dHash) algorithm, dropping >88% of redundant frames and limiting idle CPU load to <5%.

2. Accessibility-First Semantic Ingestion: Traversed native OS accessibility trees (Windows UIA, macOS AXUIElement, Linux AT-SPI) directly from RAM for pixel-perfect text extraction, utilizing on-device OCR strictly as an intelligent fallback.

3. Zero-Knowledge Cryptographic Vault & MCP: Engineered the DESKVAULT1 ChaCha20-Poly1305 AEAD container with automated PII redaction and exposed an official Model Context Protocol (MCP) server for local coding agents.

Literature Survey: Evolution of Desktop Memory Systems

Chronological Taxonomy of Lifelogging, Screen Recorders, and Ambient AI

Era / Approach	Key Reference	Major Contribution	Primary Limitations
Electro-Mechanical Indexing	Bush (1945) Memex	Conceptualized associative trails across microform desktop records.	Purely theoretical concept; completely lacked digital computing & search.
Lifelogging Systems	Gemmell & Bell (2003) MyLifeBits	Pioneered mass ingestion of digital photos, scanned docs, and audio.	Required massive manual tagging; no automated speech-to-text or NLP.
Fixed-FPS Video Loggers	Desktop Recorders (2010–2018)	Continuous 15–30 FPS screen capture and local video compression.	Massive CPU/GPU strain; generated hundreds of gigabytes per week.
Cloud Meeting SaaS	Otter.ai / Granola (2020–2023)	Cloud-based speech-to-text and automated LLM meeting summaries.	100% cloud telemetry; severe privacy/regulatory liabilities; no screen context.
Proprietary NPU Recorders	Microsoft Recall / Rewind (2024)	Periodic timer snapshots and local neural processing on NPUs.	Closed-source binaries; vendor lock-in; lack of open agent protocols (MCP).
Ambient Event-Driven Engine	Desk-AI (Present Work)	Event-driven dHash, UIA accessibility-first, local Whisper, MCP.	Requires local compute resources; platform-specific OS accessibility quirks.

Problem Statement: The Digital Context Bottleneck

Context Severance, Prompt Fatigue, and Cloud Privacy Vulnerabilities

- **The Context Severance Barrier:** Large Language Models operate as isolated, stateless oracles. They possess zero passive visibility into running terminal errors, IDE tabs, documentation windows, or meeting conversations, severing the model from the user's active computational reality.
- **Prompt Fatigue & Cognitive Friction:** Knowledge workers waste 15–20% of their workday taking screenshots, copy-pasting stack traces, and writing elaborate prompts just to give an AI agent necessary context, transforming AI into an ad-hoc chore rather than an ambient co-pilot.
- **The 'Continuous Video' Hardware Tax:** Brute-force continuous screen recording drains battery within hours, triggers thermal throttling, and fills hard drives at ~100 GB/month, rendering conventional approaches impractical on laptops.
- **The Cloud Privacy & Surveillance Chasm:** Commercial lifeloggers routinely stream unredacted screen frames and audio to remote SaaS servers, exposing enterprise trade secrets, passwords, customer data, and personal privacy to external data leaks.

THE PROMPT FATIGUE TAX

1. Active Developer Work
(IDE, Terminal, Browser)
- ▼
2. Breakpoint / Meeting
(Transient error or decision)
- ▼
3. Cognitive Friction
(Manual screenshot & formatting)
- ▼
4. Isolated LLM Prompt
(Chat box query without context)
- ▼
5. Result: Disjointed Context
(~20% Workday Wasted)

Introducing Desk-AI: The Local-First Ambient Memory Engine

Continuous, Zero-Knowledge Multimodal Context Capture for Agentic Computing

1. MULTIMODAL INGESTION

- Event-driven screen capture (WGC / ScreenCaptureKit)
- Dual-channel audio loopback (WASAPI / CoreAudio)
- Native OS accessibility trees (Windows UIA / AXUIElement)
- Silero Voice Activity Detection (VAD)

2. LOCAL ZERO-KNOWLEDGE STORE

- Single-writer SQLite Actor queue (WAL mode)
- Full-text search with BM25 ranking (FTS5)
- Automated local PII redaction & pixel blackout
- DESKVAULT1 ChaCha20-Poly1305 AEAD container

3. AGENTIC INTEROPERABILITY

- Standard Model Context Protocol (MCP) server
- Native integration with Claude, Cursor, and Codex
- Autonomous cron-scheduled agent Pipes (`pipe.md`)
- Sub-40ms semantic search via ONNX embeddings

Core Architectural Paradigm Shift:

Desk-AI fundamentally replaces the reactive 'ask and prompt' model with proactive ambient computing. Built entirely in systems-level Rust, the daemon runs silently as an unprivileged background process, recording and indexing 100% of the user's computing lifecycle without allowing a single byte of private data to leave the local machine.

By uniting high-speed accessibility extraction with local neural speech models and zero-knowledge encryption, Desk-AI transforms an ephemeral desktop into an auditable, high-dimensional memory substrate accessible to autonomous coding agents.


System Architecture: Multi-Tier Modular Topology

A Coordinated Ecosystem of Specialized Systems Intelligence in Rust

- **Multimodal Ingestion Tier:** Contains platform-specific capture drivers: Windows Graphics Capture (WGC), ScreenCaptureKit (macOS), WASAPI loopback audio, and global OS input event hooks for mouse and keyboard monitoring.
- **Processing & Neural Tier:** Executes 64-bit dHash perceptual filtering, Silero VAD voice activity segmentation, on-device Whisper Large-v3-Turbo speech transcription, speaker voiceprint clustering, and automated PII redaction.
- **Storage & Persistence Tier:** Routes all database mutations across an MPSC channel to a dedicated Single-Writer SQLite Actor operating in WAL mode, driving FTS5 BM25 ranked inverted indices with sub-15ms keyword query latencies.
- **Cryptographic Vault Tier:** DESKVAULT1 streaming encryption engine utilizing ChaCha20-Poly1305 AEAD with memory-hard Argon2id key derivation, sealing master keys directly inside OS hardware enclaves (Credential Guard / Keychain).
- **Extensibility & Agent Gateway:** High-throughput Axum REST/WebSocket server, official Model Context Protocol (MCP) server, and autonomous scheduled Markdown Pipes with deterministic token-gated OS sandboxing.
- **Dynamic Concurrency & Safety:** Orchestrated via Tokio asynchronous runtime, ensuring complete memory safety without garbage collection, zero runtime telemetry, and zero SQLITE_BUSY lock contention.

Desk-AI End-to-End Processing Pipeline

Sequential Signal Conditioning, Filtering, Local Inference, and Agent Querying



Pipeline Engineering Principles:

1. Early Discard: 88%+ of frames are dropped at Step 3 via hardware popcnt dHash before incurring expensive OCR or disk writes.
2. Accessibility Over OCR: Step 4 queries exact OS UI Automation memory nodes, achieving 100% character accuracy on code and URLs.
3. Zero-Knowledge Guarantees: Sensitive tokens are stripped and encrypted before persistence, ensuring the database is cryptographically secure at rest.

Deep Dive: Event-Driven Screen Ingestion & Perceptual dHash

Eliminating Redundant Video Capture via 64-Bit Difference Hashing

- **OS Event Tap Hooks:** Desk-AI listens to system event streams (mouse clicks, app focus switches, window repositioning, and keyboard typing pauses) rather than executing blind periodic polling. Captures occur only when the human workspace state changes.

- **Bilinear Matrix Downsampling:** Captured high-resolution display buffers (up to 4K RGB) are downsampled to a compact 9x8 grayscale matrix using bilinear filtering, eliminating high-frequency display noise, cursor blinks, and subpixel rendering artifacts.

- **Row-Wise Gradient Extraction:** Adjacent horizontal pixels in each row are compared to produce an 8x8 binary bitmap:

$\text{Bit}(x, y) = 1 \text{ if } \text{Pixel}(x, y) > \text{Pixel}(x+1, y) \text{ else } 0$

This yields a robust, orientation-sensitive 64-bit unsigned integer fingerprint (`u64`).

- **Hardware Popcount Hamming Distance:** The Hamming distance between consecutive frame dHashes is evaluated via the CPU's native `popcnt` instruction in under 1 nanosecond. If Hamming Distance < 3 bits, the frame is classified as visually identical and dropped.

- **Performance & Battery Impact:** Empirical tests confirm >88% frame reduction during coding and reading sessions, slashing storage growth from ~120 GB/month down to <15 GB/month and keeping capture CPU load under 5%.

```
// Rust: desk-ai-screen/dhash.rs

pub fn calculate_dhash(img: &GrayImage) -> u64 {
    let r = imageops::resize(
        img, 9, 8, FilterType::Nearest
    );
    let mut hash: u64 = 0;
    for y in 0..8 {
        for x in 0..8 {
            if r.get_pixel(x, y)[0] >
                r.get_pixel(x + 1, y)[0] {
                hash |= 1 << (y * 8 + x);
            }
        }
    }
    hash
}

// Fast Hamming Distance:
let diff = (h1 ^ h2).count_ones();
if diff < 3 { return Discard; }
```

Deep Dive: Dual-Stream Audio, Silero VAD & Local Whisper STT

Capturing Both Sides of Verbal Conversations Without Remote Audio Exfiltration

- **Synchronous Dual-Channel Audio Loopback:** Captures system audio output (WASAPI Loopback on Windows, CoreAudio on macOS) to record remote meeting attendees, while simultaneously recording the user's local microphone in separate 16 kHz 16-bit mono PCM buffers.
- **Silero Voice Activity Detection (VAD):** A lightweight 5-layer recurrent neural network operates on 30ms audio windows (512 samples). Continuous speech probability is tracked; if probability < 0.5, audio is dropped immediately to avoid loading speech models on silence.
- **On-Device Whisper Large-v3-Turbo:** Speech segments are transcribed entirely locally via Whisper Large-v3-Turbo executed through ONNX Runtime (`ort`) or quantized Whisper.cpp, extracting text, timestamps, and confidence scores with zero cloud API latency.
- **Acoustic Speaker Diarization & Centroids:** Computes 512-dimensional voiceprint embeddings for active audio chunks, clustering speaker identities across meeting sessions into distinct speaker centroids ('Speaker 0', 'Speaker 1') aligned with calendar events.
- **Clock Drift Resampling Compensation:** Hardware audio sample clocks drift relative to monotonic system clocks over extended captures. Desk-AI integrates `libsamplerate` automated resampling, maintaining audio-to-screen alignment within 50ms.

Deep Dive: Zero-Knowledge Security & DESKVAULT1 Encryption

Hardware-Attested Data Isolation, Argon2id Key Derivation, and Automated PII Redaction

- **Pre-Persistence PII Redaction (desk-ai-redact):** All extracted screen text and audio transcripts pass through a local Named Entity Recognition (NER) model and regex classifier. Credit cards, SSNs, passwords, and API keys are pseudonymized (`[REDACTED\_CREDIT\_CARD]`) and their bounding boxes blacked out on stored image frames.
- **The DESKVAULT1 Streaming Container:** Encrypted files are segmented into 64 KB blocks, each protected with ChaCha20-Poly1305 Authenticated Encryption with Associated Data (AEAD). Nonces are generated dynamically: `ChunkNonce = BaseNonce XOR ChunkIndex` to prevent nonce reuse.
- **Argon2id Memory-Hard Key Derivation:** User master keys are stretched using Argon2id (m=64MB, t=3, p=4), providing state-of-the-art resistance against GPU/ASIC brute-force dictionary attacks. Master keys are sealed inside OS hardware keychains (Windows Credential Guard, Apple Keychain).
- **Hardware DRM & Incognito Window Suppression:** Monitors window security attributes and process identifiers. When HDCP-protected media players (Netflix, Spotify) or browser incognito sessions enter focus, capture is immediately suspended to preserve privacy and prevent black-screen artifacts.
- **Confidential Enclave Attestation:** Integrates Tinfoil confidential computing (`tinfoil-rs`), validating that cryptographic key operations execute within hardware-attested Trusted Execution Environments (TEEs).

Backend Data Architecture: SQLite WAL & FTS5 Indexing


High-Throughput Concurrency, Normalized Relational Schemas, and Sub-40ms Search

Table Name	Primary Keys / Foreign Keys	Indexed Attributes	Functional Description & Semantics
frames	id (INTEGER PK)	timestamp, app_name, browser_url	Visual capture snapshots, active window titles, URLs, and 64-bit dHash fingerprints.
elements	id (INTEGER PK), frame_id (FK)	frame_id, role, activity_summary	Normalized accessibility tree UI nodes (buttons, text fields, URLs, bounding boxes).
audio_transcriptions	id (INTEGER PK), chunk_id (FK)	speaker_id, timestamp, engine	Whisper speech transcripts, detected languages, speaker IDs, and confidence scores.
meetings	id (INTEGER PK)	start_time, end_time, app_name	High-level meeting sessions mapped from active collaboration software (Zoom/Teams).
frames_full_text	Virtual FTS5 Table	content, window_name, app_name	Inverted BM25 ranked full-text search index across all visual and acoustic tokens.

- **Single-Writer MPSC Actor Pattern:** All capture threads send write operations through a crossbeam MPSC channel to a dedicated background writer, completely eliminating SQLITE_BUSY lock contention.
- **Write-Ahead Logging (WAL) Concurrency:** PRAGMA journal_mode = WAL decouples readers from writers, permitting simultaneous, non-blocking full-text search queries from MCP agents while high-frequency capture commits.

User Interface: Desktop Application & Control Center

Escher Monochrome Interface with Instant Chronological Scrubber (Tauri v2 + Next.js)



- **Interactive DVR Timeline Scrubber:** Users can scrub backwards seamlessly across their entire workday with sub-second granularity, inspecting visual states, active tabs, and meeting transcripts.
- **Multimodal Synchronized Inspection:** Displays synchronized video frame snapshots, OCR bounding boxes, and time-aligned speech transcripts in a single unified view.
- **Native App & Connector Registry:** One-click local connectors for Google Calendar, Apple Calendar, Gmail, Obsidian, Notion, and Linear for automatic activity correlation.
- **Instant Privacy Toggles & Ignore Lists:** Features a 1-click recording pause toggle, dynamic application deny-lists (e.g. 1Password), and customizable disk storage auto-pruning rules.

Agent Interoperability: Model Context Protocol (MCP) & Pipes

Bridging Desktop Memory Context to Autonomous AI Coding Agents

- **Official Model Context Protocol (MCP) Server:** Implemented in Rust (``desk-ai-mcp``), exposing standardized JSON-RPC tools over bi-directional STDIO and HTTP Server-Sent Events (SSE). Seamlessly connects with Claude Desktop, Cursor AI, and Continue.dev.


- **Specialized Agent Context Tools:** Exposes high-level tools:

- ``search_screen_history(q, time_range)``
- ``get_recent_context(window_minutes)``
- ``query_meeting_transcripts(meeting_id)``

Enables coding agents to autonomously inspect past compiler errors, terminal commands, or discussions.

- **Autonomous Scheduled Pipes Engine:** Automated workflows defined in Markdown files (``pipe.md``) with YAML frontmatter. Executed on cron schedules, pipes run agents that summarize meetings into Obsidian or track build failures.

- **Deterministic OS Security Sandboxing:** Pipes and external agent integrations are strictly sandboxed. Frontmatter rules (``allow-apps``, ``deny-apps``, ``allow-raw-sql: false``) are verified via cryptographic token gating before any query executes.



Evolution: Mid-Term Prototype vs. End-Term System

Architectural Milestones from Experimental Screen Recorder to Production Engine

- **Phase 1: Screen Ingestion Evolution:** Mid-Term: Periodic timer snapshot every 3 seconds; consumed 18–25% CPU and generated thousands of duplicate frames.

End-Term: OS event taps (clicks/keys) + 64-bit dHash perceptual filtering; drops >88% of redundant frames, slashing CPU to <5%.

- **Phase 2: Text Ingestion Evolution:** Mid-Term: Relied solely on local Tesseract/Windows Media OCR; slow (>350ms per frame), prone to syntax errors on code and URLs.

End-Term: Native OS Accessibility Tree traversal (UIA COM / AXUIElement); extracts exact text from RAM with OCR strictly as fallback.

- **Phase 3: Acoustic Pipeline Evolution:** Mid-Term: Naive continuous audio recording; transcribing long periods of empty silence wasted CPU cycles.

End-Term: 5-layer Silero VAD segmentation + local Whisper Large-v3-Turbo + acoustic speaker diarization clustering.

- **Phase 4: Persistence & Storage Evolution:** Mid-Term: Direct multi-threaded SQLite connections prone to SQLITE_BUSY database locks and unencrypted disk storage.

End-Term: Single-Writer Actor queue in WAL mode + FTS5 BM25 search + DESKVAULT1 ChaCha20-Poly1305 streaming encryption.

- **Phase 5: Agent Extensibility Evolution:** Mid-Term: Basic internal keyword search accessible only through desktop UI.

End-Term: Full Model Context Protocol (MCP) server + autonomous scheduled Markdown Pipes with token sandboxing.

Experimental Benchmarks & Performance Evaluation

Empirical Profiling Across 8-Hour Continuous Workloads on Intel Core i7 & Apple M2

Operational Workload State	CPU Usage (%)	RAM Footprint	Disk Write Rate	Thermal Impact
Idle (Screen static, room silence)	0.8% - 1.5%	480 MB	< 2 MB / hour	Ambient / Negligible
Active Web Browsing & Code Editing	3.2% - 5.8%	720 MB	18 MB / hour	Nominal (+0.5°C)
Video Conference Call (WASAPI + VAD)	6.5% - 9.4%	1,150 MB	45 MB / hour	Mild (+1.5°C)
Peak Whisper Transcription (Local Batch)	18.2% (burst)	2,200 MB	85 MB / hour	Transient Fan Ramp

- **Query Retrieval Latency (1.2M Rows):** Exact B-Tree match: 4.2 ms (p50) / 12.8 ms (p95) | FTS5 BM25 keyword search: 14.5 ms (p50) / 38.2 ms (p95) | Dense vector cosine similarity: 42.0 ms (p50) / 89.5 ms (p95).
- **Validation Verdict:** Desk-AI satisfies all non-functional requirements: continuous 24/7 background operation consumes under 6% average CPU, less than 2 GB RAM, and under 15 GB of compressed disk storage per month.

Why Commercial Solutions Fail for Desktop Context

Cloud Recorders and Closed Systems Fundamentally Compromise Enterprise Security and Performance

Context Severance & Prompt Fatigue

Standard commercial LLMs cannot observe active desktop state. Users are taxed with taking manual screenshots, copying terminal stack traces, and formatting prompts.

Desk-AI Advantage: Provides continuous ambient episodic memory directly to AI coding agents via standard MCP protocols.

Cloud Exfiltration & Surveillance Risk

Commercial meeting and screen recorders (Otter, Granola, Recall) upload audio and unredacted screens to external SaaS servers, violating GDPR, HIPAA, and IP secrecy.

Desk-AI Advantage: Operates 100% on-device with zero telemetry and chunked DESKVAULT1 ChaCha20-Poly1305 encryption.

Resource Bloat & Severe Battery Drain

Brute-force fixed-framerate screen video capture drains mobile laptop batteries within 2 hours, triggers thermal throttling, and fills >100 GB of storage monthly.

Desk-AI Advantage: Event-driven dHash filtering limits CPU load to <5% and monthly disk growth to <15 GB.

Proprietary Lock-in & Closed Silos

Existing solutions trap activity data in proprietary, closed-source binary formats, preventing external tooling integration or custom automation pipelines.

Desk-AI Advantage: Built on standard SQLite WAL, open FTS5 virtual tables, and autonomous Markdown Pipes.

Verification, Testing & Quality Assurance

Rigorous Systems Validation for Continuous 24/7 Background Daemon Reliability

Test ID	Subsystem	Input Test Stimulus	Expected Specification Behavior	Status
TC-01	Screen Capture	Static desktop view with 0 user input for 60s	dHash matches; 0 redundant frames committed to disk	PASSED
TC-02	Accessibility	Browser tab switch from IDE to documentation	UIA tree serializes new URL & title in <30ms	PASSED
TC-03	Audio VAD	30-second silent background audio stream	Silero VAD flags prob < 0.2; 0 Whisper STT calls	PASSED
TC-04	PII Redaction	Display showing mock credit card: 4532 0156 ...	Text pseudonymized; frame bounding box blacked	PASSED
TC-05	Crash Recovery	Simulated SIGKILL during active SQLite WAL write	Automatic WAL replay clean; 0 database corruption	PASSED

- **Failpoint Virtual File System (VFS) Testing:** Injected 5,000 simulated kernel write panics and abrupt power loss faults during SQLite checkpoints; achieved 100% automatic recovery with zero data loss.
- **End-to-End Headless Daemon Test Suite:** Simulated synthetic mouse and keyboard workflows; validated that captured events appear in MCP search outputs within 1.5 seconds.

References

Formal Academic Literature & Architectural Specifications

- [1] V. Bush, “As We May Think,” The Atlantic Monthly, vol. 176, no. 1, pp. 101–108, Jul. 1945.
- [2] J. Gemmell, G. Bell, R. Lueder, S. Drucker, and C. Wong, “MyLifeBits: Fulfilling the Memex Vision,” ACM Computers in Entertainment, vol. 1, no. 1, pp. 1–12, Oct. 2003.
- [3] A. Radford, J. W. Kim, T. Xu, G. Brockman, C. McLeavey, and I. Sutskever, “Robust Speech Recognition via Large-Scale Weak Supervision (Whisper),” in Proc. Int. Conf. on Machine Learning (ICML), pp. 28492–28518, 2023.
- [4] Silero Team, “Silero VAD: Pre-trained Enterprise-Grade Voice Activity Detector,” GitHub Repository, 2021.
- [5] D. R. Hipp, “SQLite: An In-Process Database Engine with Full-Text Search (FTS5),” SQLite Consortium Documentation, 2024.
- [6] A. Biryukov, D. Dinu, and D. Khovratovich, “Argon2: New Generation of Memory-Hard Password Hashing Functions,” in Proc. IEEE European Symposium on Security and Privacy (EuroS&P), pp. 292–302, 2016.
- [7] D. J. Bernstein, “ChaCha, a Variant of Salsa20,” in SASC 2008 - The State of the Art of Stream Ciphers, pp. 273–278, 2008.
- [8] Anthropic PBC, “Model Context Protocol (MCP) Specification,” Architecture Documentation, <https://modelcontextprotocol.io>, 2024.